

Serial Optimization Tips

HPC BOOTCAMP

- Optimization, bag-of-tasks, MPI, OpenMP
- Week of June 14
- Free!
- Goodies provided for morning and afternoon breaks
- Python is taught except for OpenMP

In Nearly All Cases The Tradeoff Is:

- Memory versus speed
 - More memory usage=more speed UNTIL you use enough memory that you slow down (or crash) the entire system
 - More memory usage=fewer options for running (less of a problem nowadays)
 - But if you do use enough memory to slow down the system it will be MUCH slower.
 - Please try to avoid 4000x4000 arrays!

For All Languages:

- Do not recompute
 - We see many codes where programmers recompute the same quantity over and over
 - If it's a scalar like $\pi/180.0$ this isn't horrible, but it's not great either, and it's not necessary! Compute it first and store it.
 - Do not recompute *anything* in a big loop that you can precompute and store.
 - Do not recompute information already accumulated even if you must introduce a new array (unless it's huge).

Loop Efficiency

- Do not recompute array accesses. Make a variable

```
do j=1,10
  do i=1,10
    Something=r(j)*sin(z(i-j+1))*x(i,j)
  Enddo
enddo
```

Better

```
do j=1,10
  Rj=r(j)
  j1=1-j
  do i=1,10
    Something=Rj*sin(z(i+j1))*x(i,j)
  enddo
enddo
```

More Loop Efficiency

- Avoid conditionals in inner loops
- Sometimes it's best to go through large multidimensional arrays as if they were linear (which they are in memory).
- Or just use an array operation and trust the compiler/interpreter

Reduce Math Functions

- Fast:
 - Add/subtract/multiply (fma: fused multiply add is usually a hardware instruction also, i.e. $a = a + b * c$)
- Slow:
 - Division
- Really slow:
 - Math functions especially power (to noninteger in Fortran, to anything in other languages) and square root and such.

Fortran

- Make sure loops are in column-major order (right to left)

```
do j=1,10000
  do i=1,10000
    z(i,j)=a(i,j)+b
  enddo
enddo
```

- This is for *stride 1 memory access* and *cache efficiency*
- Opposite for row-major order languages (C/C++ etc.)

Python

- Avoid dynamically resized lists, or anything that changes dynamically during the run after it is initialized
 - OK to allocate/size at runtime, but once the size is fixed, allocate it and keep it the same size
 - NumPy array versus list
 - Frozenset versus set
- Avoid `for` loops. Use almost anything else:
 - List comprehension
 - Iterator

Python Example

```
for i in range(1,n-1):  
    for j in range(1,n-1):  
        u[i,j]=0.25*(u[i+1,j]+u[i-1,j]+u[i,j\  
+1]+u[i,j-1])
```

- Can be replaced by
- $u[1:-1, 1:-1] = 0.25 * (u[2:, 1:-1] + u[-2:1:-1] + \backslash$
 $u[1:-1, 2:] + u[1:-1, -2:])$
- This can result in a speedup of up to 250x.

Python Tips and Tricks

- Don't concatenate strings, use the join method instead (or use print to format a string)
- Use map rather than a loop to apply a function to a sequence
- <http://wiki.python.org/moin/PythonSpeed/PerformanceTips>